

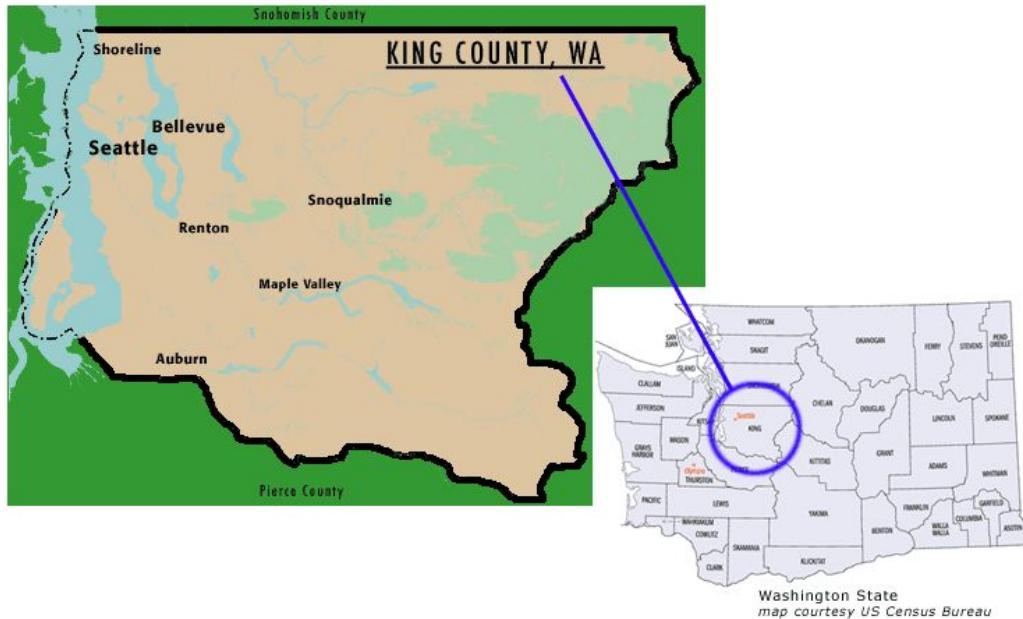
House Pricing Prediction

Machine Learning –
Supervised Learning



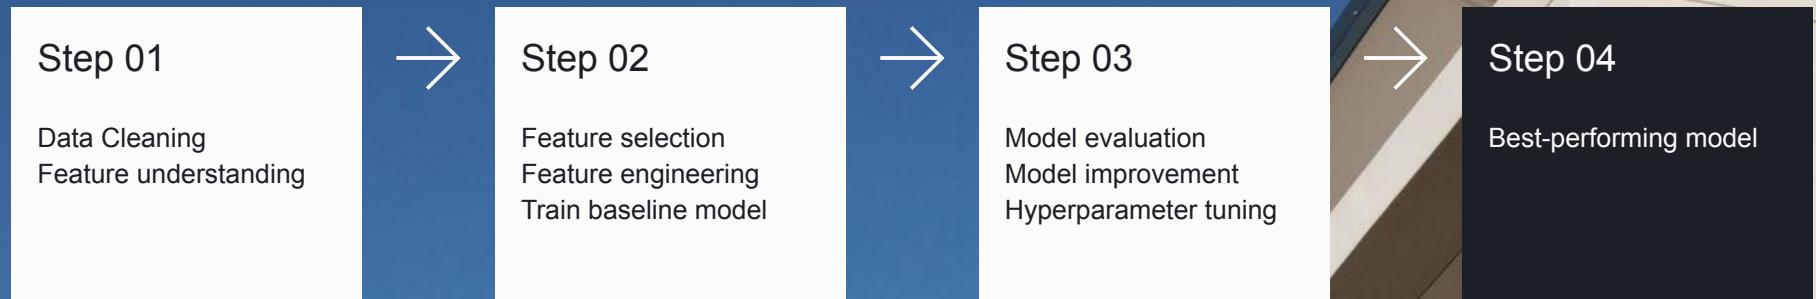
Hayley
Mohamad
Mara

Dataset House sale prices in King County (May 2014–May 2015)



- Our goal is to predict house price
- Why interesting:
 - real-world dataset
 - interpretable features
 - mix of categorical/numerical
 - 18 features: size, location, the number of rooms etc

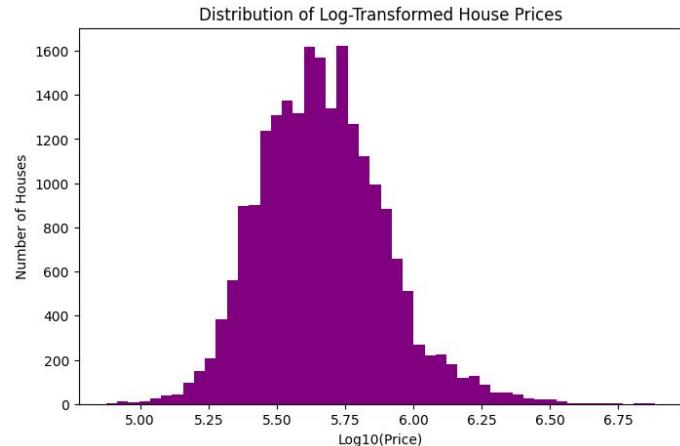
Project roadmap



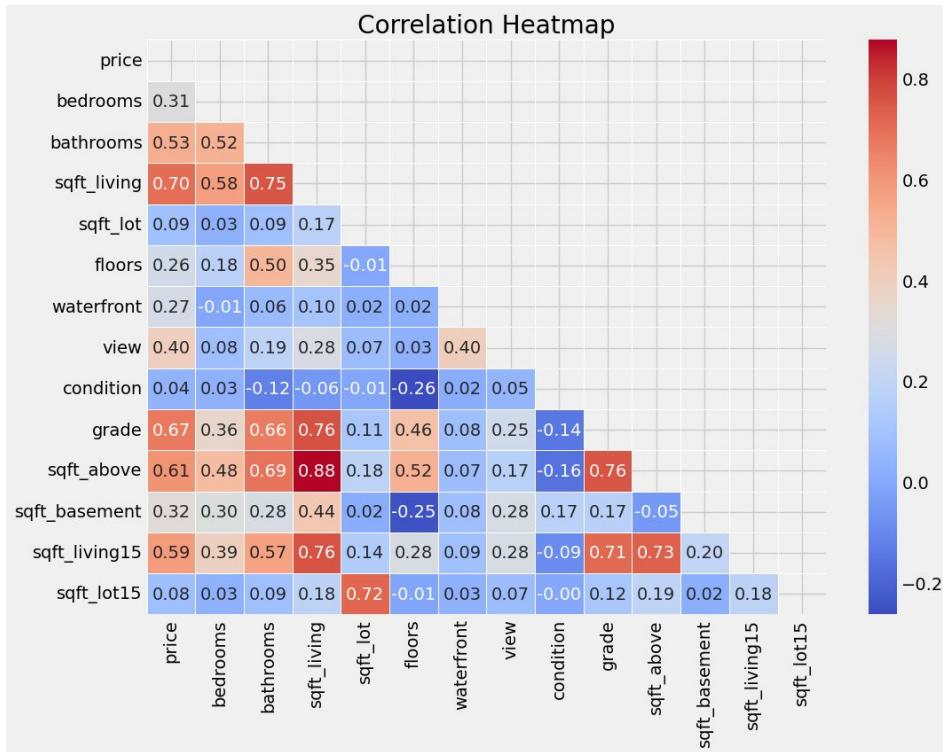
Data Understanding & Cleaning

Data Cleaning Steps:

1. no missing values
2. date -> datetime format
3. removal impossible values
 - a. bed/bath = 0
 - b. sale_year < yr_built



Feature Understanding



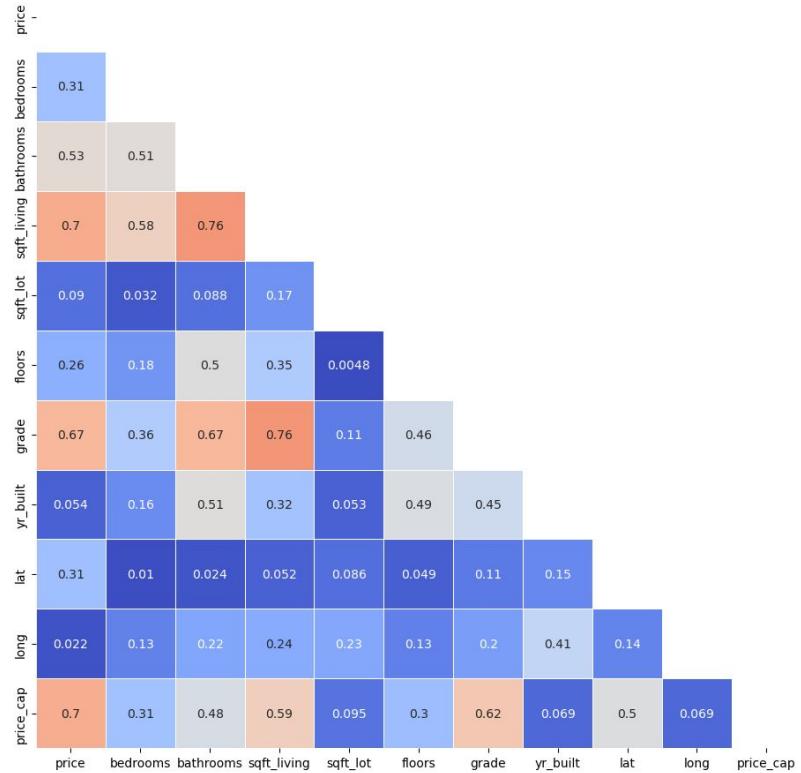
Correlation Heatmap of all features

- Highest correlation (0.88) between *sqft_living* and *sqft_above*
- Strong correlation between *sqft_living/grade* and *price*

Feature Selection

Feature	Reason to drop
ID	unique values
waterfront	only very small dataset
view	only very few unique values
ZIP	lat & long as area features
yr_renovated	we have the grade
condition	grade & yr_build
sqft_living15 & sqft_lot15	sqft_living & sqft_lot recorded in 2015
sqft_above & sqft_basement	we have sqft_living
date	only one year of data

Correlation matrix after feature selection

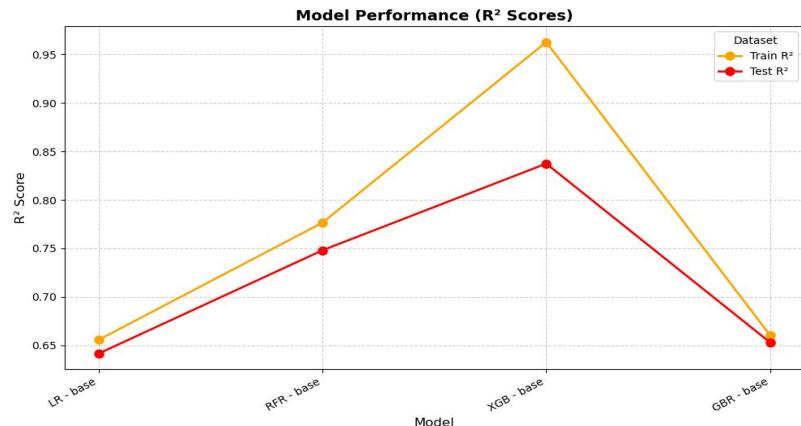
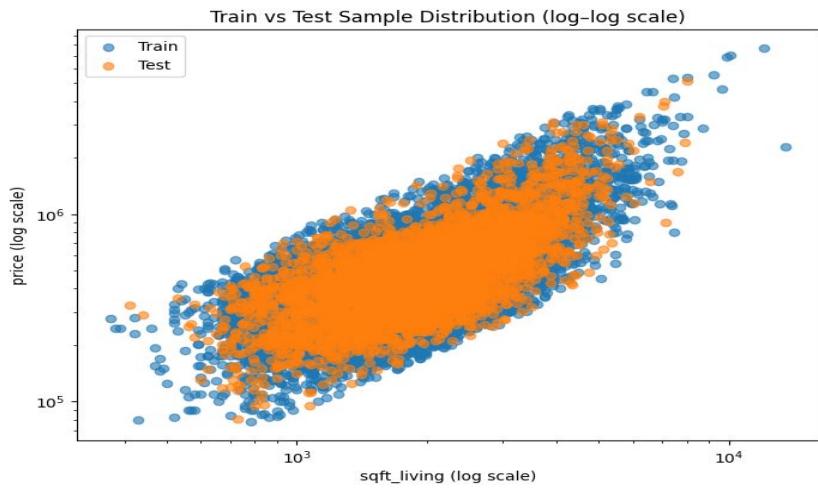


correlation of selected features with target (price)

- Relatively strong correlation with sqft_living(0.7), grade (0.67)
- Moderate with structural features (bathrooms(0.53), bedrooms(0.3), floors)
- Moderate or weak correlation location(lat(0.31), long(0.022))

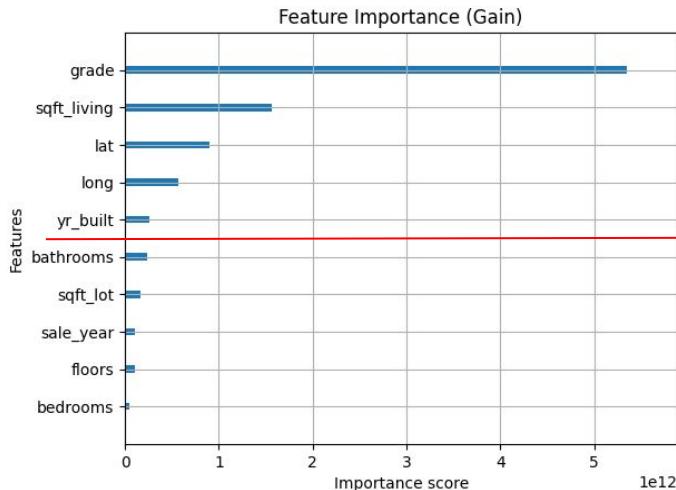
Baseline Model + Comparison

- Train/Test split 80/20, target: price
- 4 regressors:
 - Linear Regression
 - Random Forest
 - XGBoost
 - Gradient Boost
- Models assessed
- Assessment: R-Squared value.



Iteration Features & Hyperparameters – XGB

- Capping price (upper Q3)
- Feature engineering ($\text{age} = \text{sale_year} - \text{yr_built}$)
- Dropping low-impact features
- Adding view/waterfront back
- Standardization (sqft_living , year)
- Randomized Search CV for Hyperparameters



Feature Importance:

- Gain Importance: measures contribution of feature to model's accuracy.
- Quantifies average reduction in loss when feature is used in split.

Iteration Performance Features – XGB



- Trade-off between training and test performance.
- Possible overfitting for the base model.
- Adding price cap, i.e removing outliers → significant improvement.
- Model based on top features according to gain had worse training performance.
- Hyperparameter testing brought train & test closer.

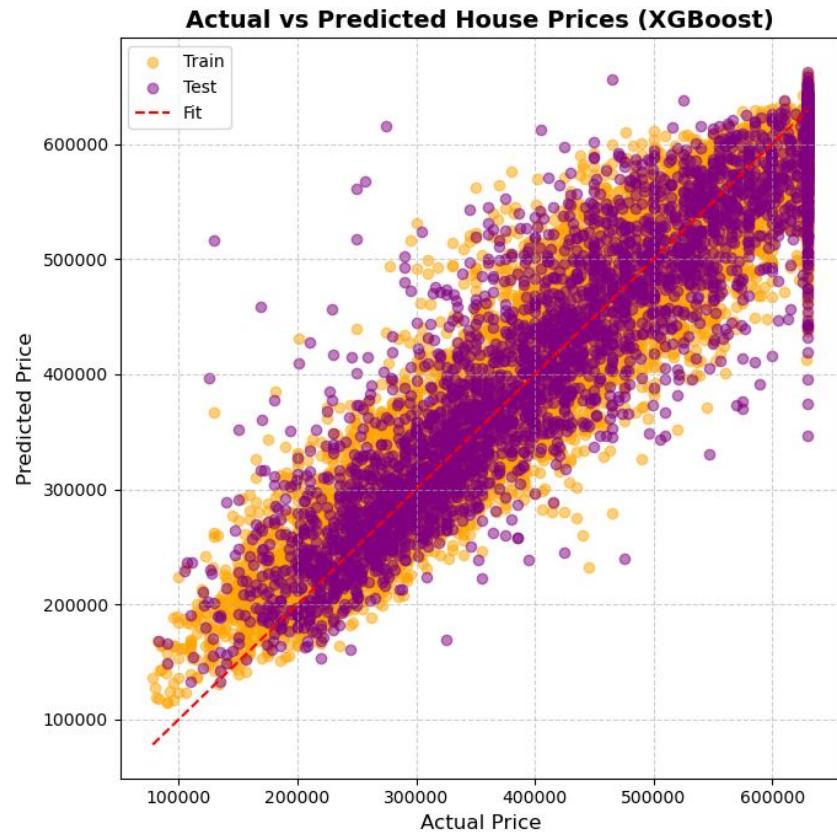
Final Model Performance

Train MAE: 26625.98
Test MAE: 35764.96
Train R²: 0.94
Test R²: 0.88

Best Parameters Found (XGBoost)	
Parameter	Value
subsample	0.6
reg_lambda	1
reg_alpha	0
n_estimators	300
max_depth	7
learning_rate	0.05
gamma	0.1
colsample_bytree	0.6

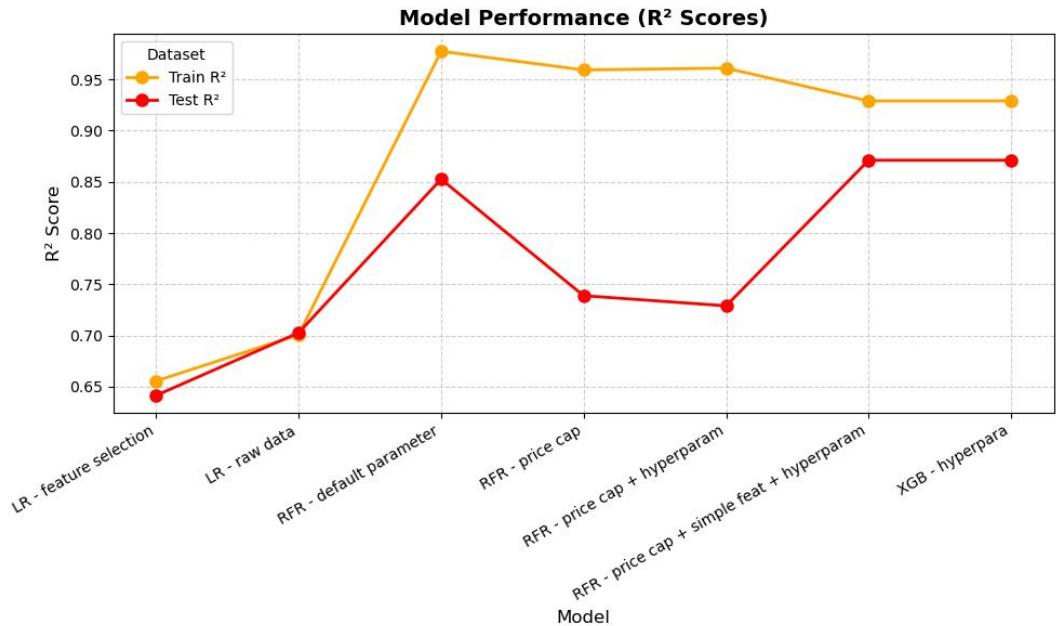
final reduced dataset + price cap:

sqft_living	grade	yr_built	lat	long
2380	8	2010	47.7170	-122.020
3190	10	1999	47.5115	-122.246
1730	8	1994	47.2621	-122.308
1870	7	1977	47.1985	-122.001
2850	7	1959	47.4859	-122.205



Learnings

- Early aggressive feature dropping *hurt* performance — some columns that seemed redundant added signal.
- Model selection (try early, tune later) is more efficient than heavy pre-filtering.
- XGBoost slightly best overall, but Random Forest nearly matched it — showing simplicity can win.
- Probably better to use $\log(\text{price})$ than $\text{capQ3}(\text{price})$.





Thank You!
