
House price prediction

Supervised machine learning

Hayley Hyejeong Lee



Project Goal

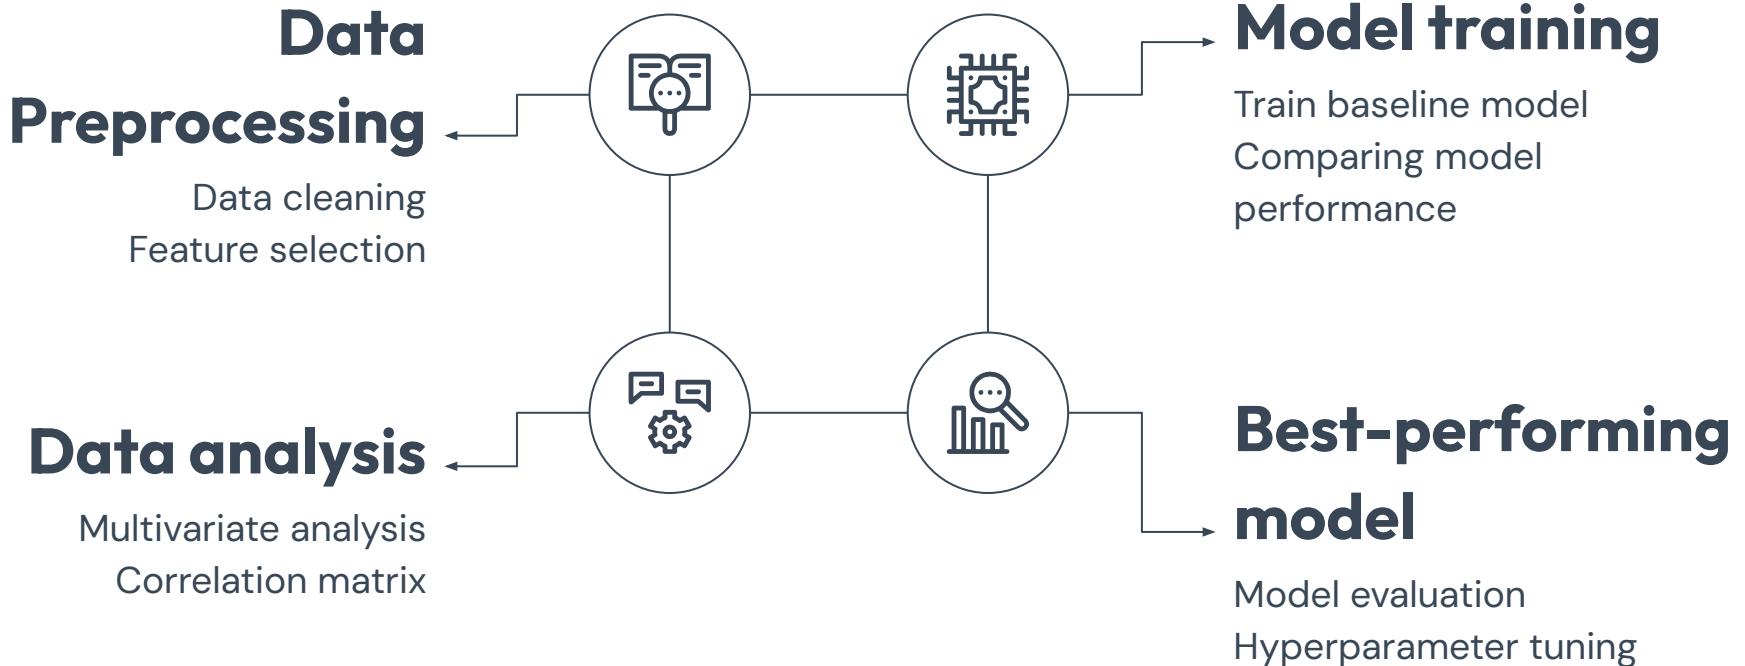
To predict the house price using
different regression models

To find the best-performing model

Dataset
House sale prices
in King County
(May 2014–May 2015)

18 features
(both numerical and categorical features)

Project Workflow



Data preprocessing

Data cleaning

- Handling missing values
- Check duplicates rows
- Drop the columns where bedrooms and bathrooms are 0 & sale_year<yr_built
- Date to datetime (dtype)
- Convert to house_age from sale_year and yr_built

Data analysis

- Univariate analysis on target variable: price
- Bivariate analysis with other features (grade, sqft_living, lat, long etc)
- Correlation matrix for numerical columns to check multicollinearity

Univariate analysis : Price

● Descriptive Statistics

Mean: 5,400,881

Min: 75,000

Max: 7,700,000

● Distribution

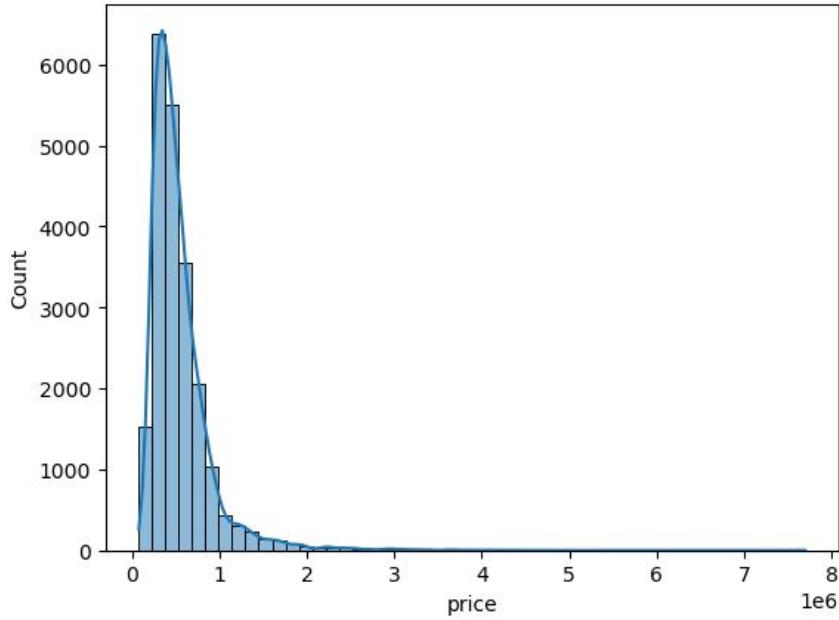
Skewness (4.02)

-> highly right-skewed

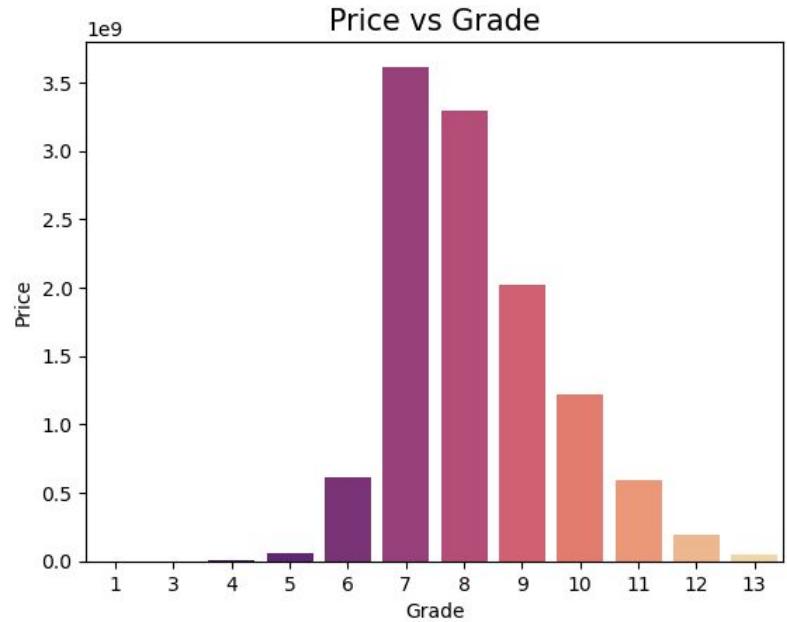
Kurtosis(34.58)

-> heavily tailed

-> contain extreme outliers (may be capping, log transformation)

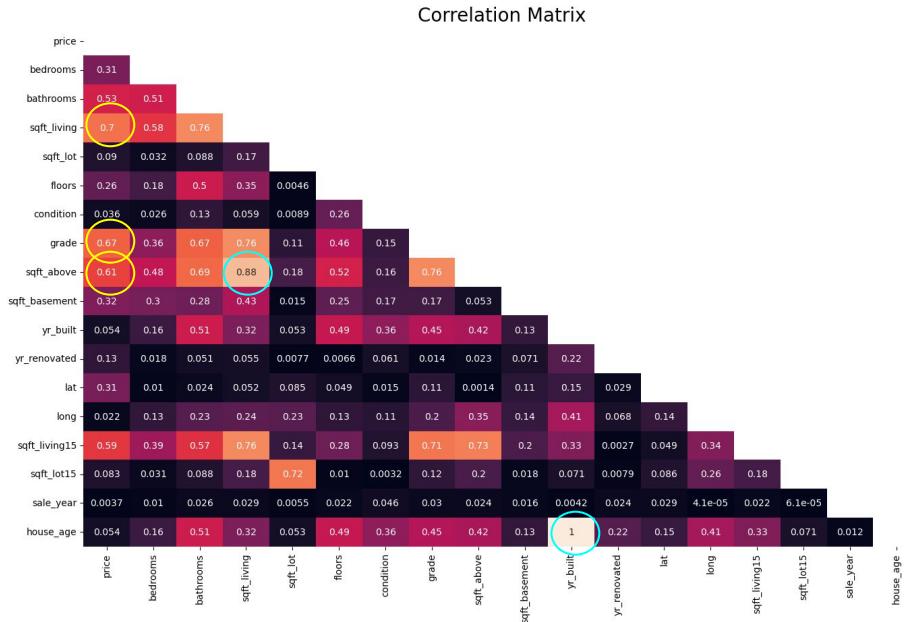


Bivariate analysis



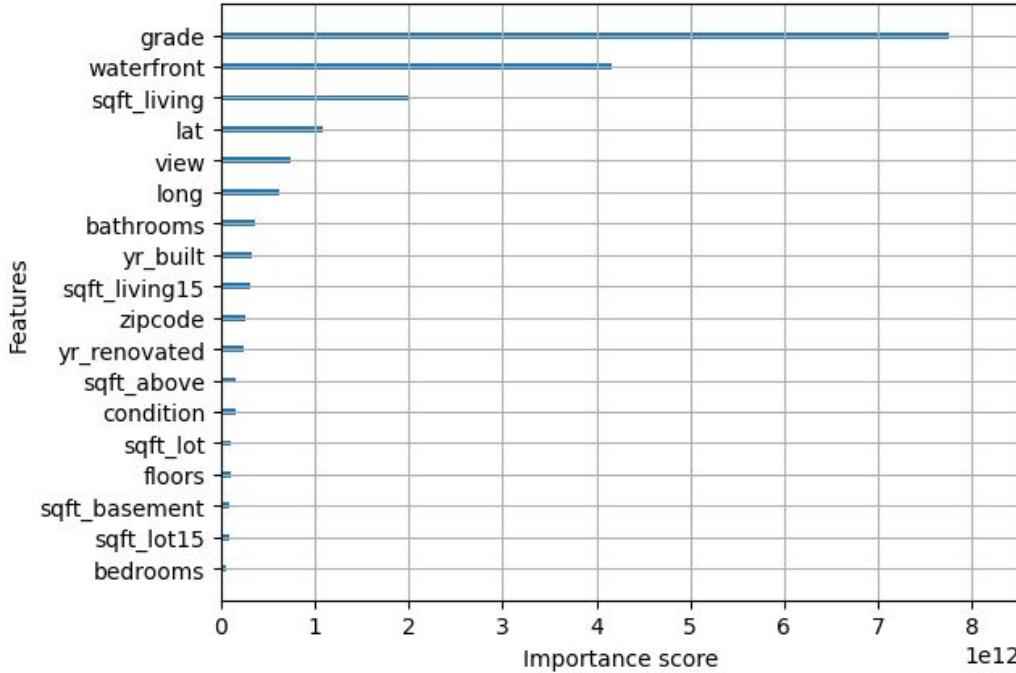
Correlation matrix

Features	Correlation score (with price)
sqft_living	0.702
grade	0.668
sqft_above	0.605
sqft_living15	0.585
bathrooms	0.526



Strong correlation (light blue circle)
 Between house_age and yr_built (1)
 Between sqft_lliving and sqft_above (0.88)

Feature importance



Top 10 features	
grade	
waterfront	
sqft_living	
lat	
long	
view	
house_age	
zipcode	
sqft_living15	

Grade,
waterfront,
sqft_living
explain about
70 % of the
total model
importance
(Xgboost)

Interaction: selected features

Data columns (total 19 columns):			
#	Column	Non-Null Count	Dtype
0	price	21585	non-null float64
1	bedrooms	21585	non-null int64
2	bathrooms	21585	non-null float64
3	sqft_living	21585	non-null int64
4	sqft_lot	21585	non-null int64
5	floors	21585	non-null float64
6	waterfront	21585	non-null int64
7	view	21585	non-null int64
8	condition	21585	non-null int64
9	grade	21585	non-null int64
10	sqft_basement	21585	non-null int64
11	yr_renovated	21585	non-null int64
12	zipcode	21585	non-null int64
13	lat	21585	non-null float64
14	long	21585	non-null float64
15	sqft_living15	21585	non-null int64
16	sqft_lot15	21585	non-null int64
17	sale_year	21585	non-null int32
18	house_age	21585	non-null int64

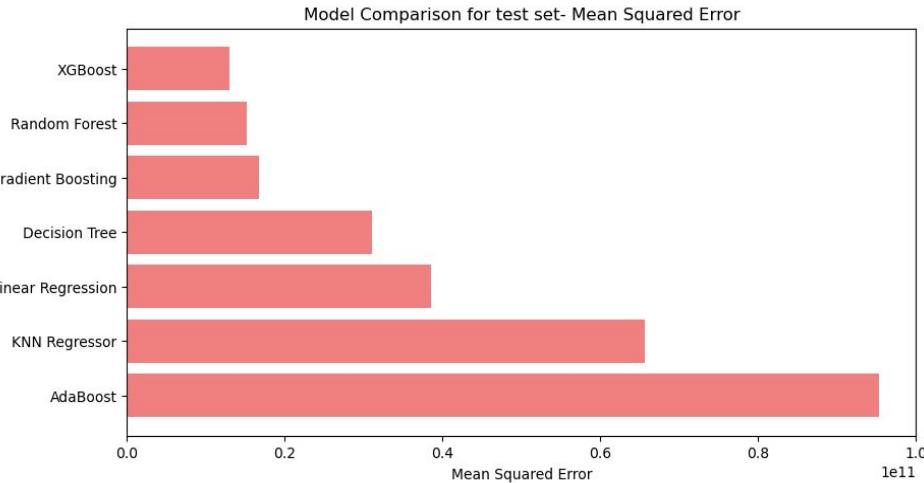
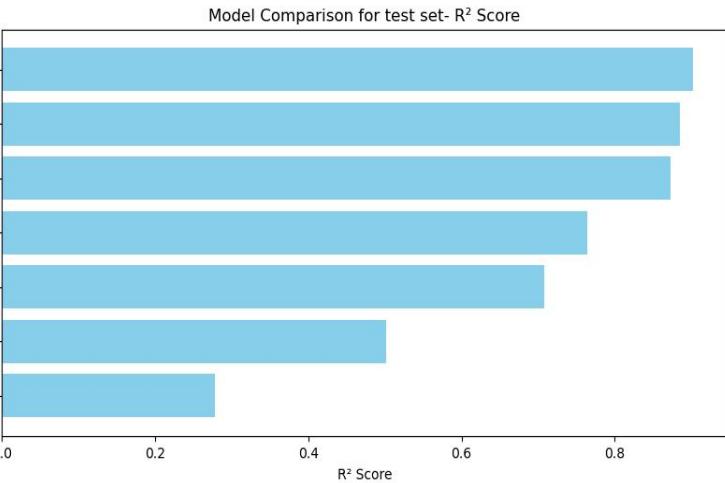
Four columns were removed
= id, date, yr_built, sqft_above

Model Comparison (Train vs Test):

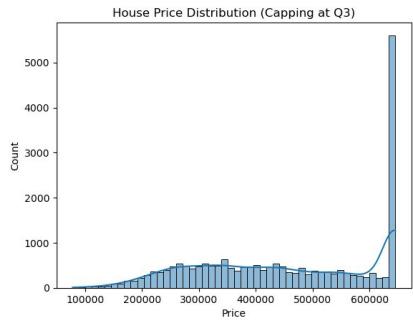
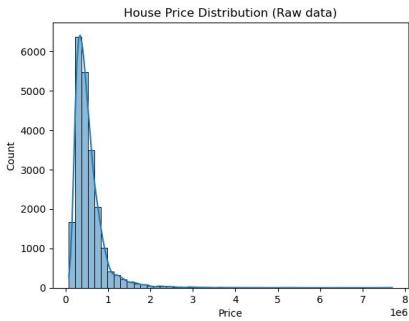
	Model	Train_MSE	Test_MSE	Train_R2	Test_R2
6	XGBoost	2.962096e+09	1.296811e+10	0.978149	0.901710
3	Random Forest	2.401300e+09	1.521478e+10	0.982286	0.884682
5	Gradient Boosting	1.319050e+10	1.671492e+10	0.902697	0.873312
2	Decision Tree	9.108929e+06	3.104438e+10	0.999933	0.764704
0	Linear Regression	4.064180e+10	3.848296e+10	0.700195	0.708325
1	KNN Regressor	4.543097e+10	6.568779e+10	0.664866	0.502130
4	AdaBoost	9.334322e+10	9.526367e+10	0.311429	0.277965

XGBoost showed the highest performance

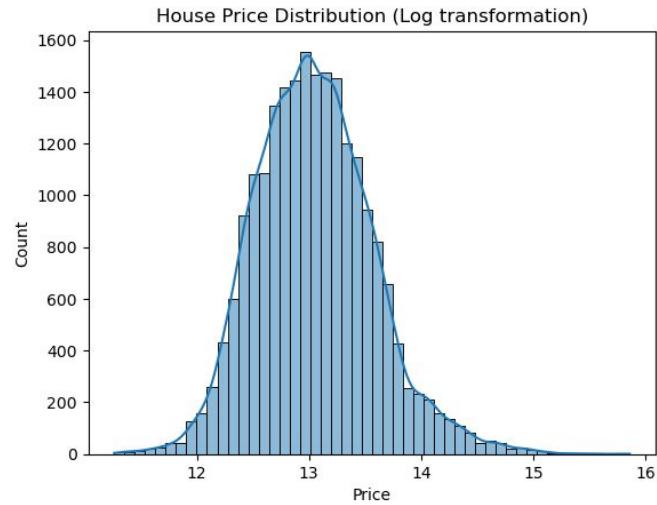
Model evaluation



Feature engineering: price



Price capping above Q3



Log transformation

Model improvement



Standardization

No impact on model performance

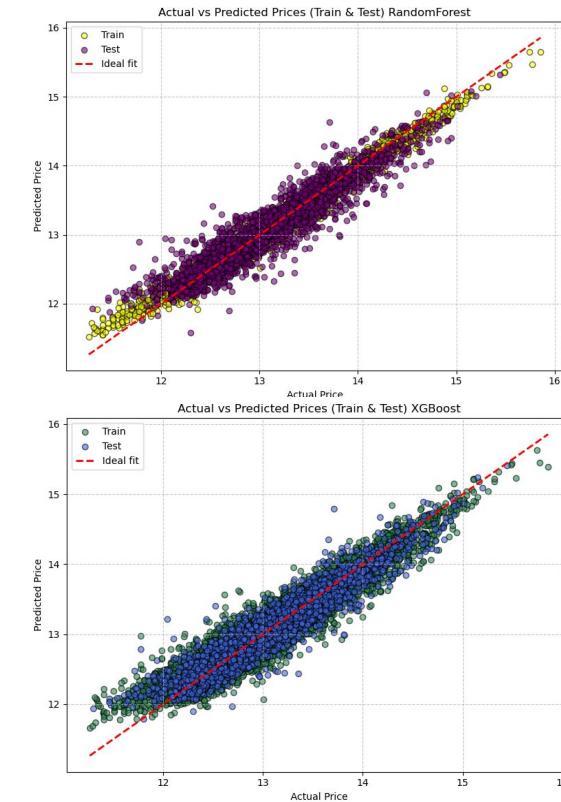
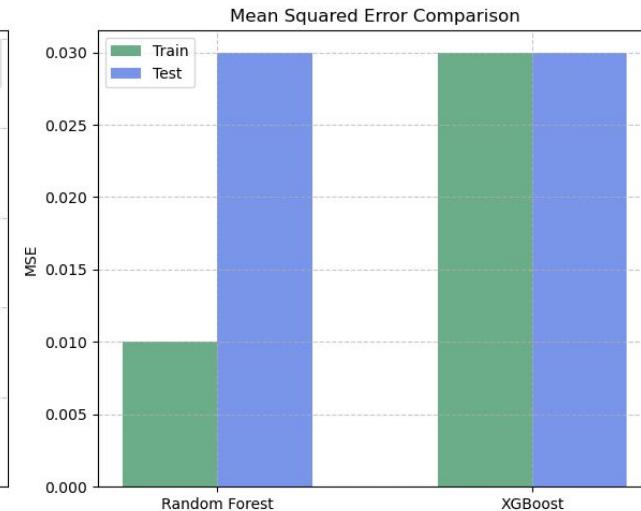
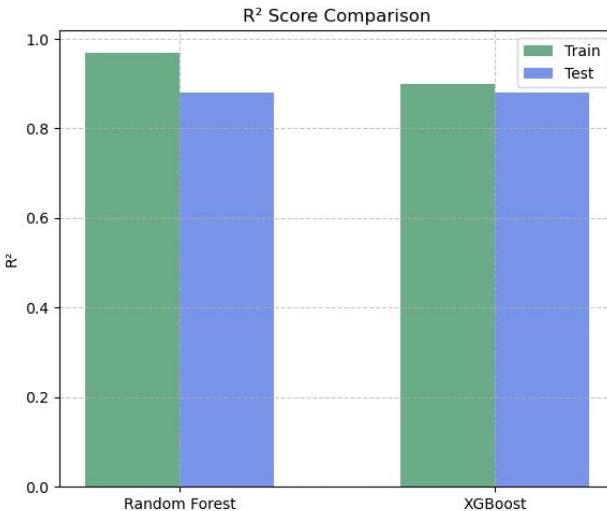
Scaled feature:
sqft_living, sqft_log, sqft_living15
sqft_lot15, sqft_basement



Hyperparameter tuning

RandomizedSearchCV
using XGBoost and Random Forest models

RandomForest vs XGBoost



Summary

XGBoost fits the testing data better (higher R2 score: 0.92)
RandomForest shows slightly more overfitting (gap between train and test, about 10%)
MSE for testing data is identical for both models.

Overall, XGBoost is more balanced, generalized better than RandomForest.



Thanks!

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)

